
Measuring Training×Refinement Interaction in Code Generation: A Factorial Framework with Mutual Information Probes

Anonymous Authors¹

Abstract

1. Introduction

2. Related Work

3. Methodology

4. Experimental Setup

5. Results

6. Discussion

7. Conclusion

References

Belghazi, M. I., Barber, D., Rajeswar, S., Oord, A. v. d., Vincent, P., and Bengio, Y. Mutual information neural estimation. In *International Conference on Machine Learning*, pp. 531–540, 2018.

Chen, o. Rethinking fine-tuning when scaling test-time compute: Limiting confidence improves mathematical reasoning. *arXiv preprint*, 2025.

Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.

Le, H., Wang, Y., Gotmare, A. D., Savarese, S., and Hoi, S. C. CodeRL: Mastering code generation through pretrained models and deep reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 35, pp. 21314–21328, 2022.

Li, D., Cao, S., et al. S*: Test time scaling for code generation. *arXiv preprint arXiv:2502.14382*, 2025.

Liu, J., Xia, C. S., Wang, Y., and Zhang, L. Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems*, volume 36, 2024.

Ma, Y., Li, B., et al. Thinking longer, not larger: Enhancing software engineering agents via scaling test-time compute. *arXiv preprint arXiv:2503.23803*, 2025.

Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhume, S., Yang, Y., et al. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

Shojaee, P., Jain, A., Tipirneni, S., and Reddy, C. K. Execution-based code generation using deep reinforcement learning. In *International Conference on Learning Representations*, 2023.

Wang, Y., Le, H., Gotmare, A. D., Bui, N. D., Li, J., and Hoi, S. C. CodeT5+: Open code large language models for code understanding and generation. *arXiv preprint arXiv:2305.07922*, 2023.

Yu, Z., Tao, Y., et al. B-Coder: Value-based deep reinforcement learning for program synthesis. In *International Conference on Learning Representations*, 2023.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.