

# Bidirectional Alignment: Teaching Language Models to Follow Any Constraint

Anonymous Authors<sup>1</sup>

## Abstract

Current RLHF optimizes language models for helpfulness alone, neglecting the *Human*  $\rightarrow$  *AI* direction of alignment: controllability. This work introduces bidirectional alignment, augmenting standard helpfulness rewards with IFEval-derived controllability signals via the combined reward  $R = \alpha \cdot R_{\text{helpfulness}} + \beta \cdot R_{\text{IFEval}}$ . Discrete constraint checks are converted into differentiable rewards using sigmoid soft thresholds.

Proof-of-concept experiments on Llama-3-8B-Instruct validate three predictions: (1) bidirectional models achieve +3.4 percentage points (pp) held-out IFEval strict accuracy over helpfulness-only baselines; (2) at  $\alpha = 0.8$ , models retain 96.4% of baseline AlpacaEval performance; (3) explicit constraint training correlates with implicit safety improvements, with +2.35pp TruthfulQA MC1 and +1.54pp BBQ improvements over baseline maxima.

The correlation between IFEval gains and safety gains ( $r=0.85-0.94$ ) suggests that explicit constraint training may build general constraint-following capacity. This work provides an initial empirical test of the bidirectional alignment framework proposed by Sun et al. (2024), introduces IFEval as a differentiable training signal, and characterizes the helpfulness-controllability trade-off at proof-of-concept scale.

## 1. Introduction

Reinforcement Learning from Human Feedback (RLHF) has become the dominant paradigm for align-

ing language models with human preferences (Ouyang et al., 2022). Current approaches optimize a single objective: helpfulness as judged by human annotators or proxy reward models. This unidirectional alignment (*AI*  $\rightarrow$  *Human*) has produced capable assistants, yet models follow instructions inconsistently, particularly explicit constraints on format, length, and structure (Zhou et al., 2023a).

Sun et al. (2024) propose a bidirectional alignment framework distinguishing two directions: *AI*  $\rightarrow$  *Human* (helpfulness, harmlessness) and *Human*  $\rightarrow$  *AI* (controllability, steerability). While the theoretical framework is compelling, no existing work implements bidirectional training signals. IFEval (Zhou et al., 2023a) measures instruction-following for evaluation; AlpacaEval (Dubois et al., 2024) measures helpfulness—but these metrics remain siloed for post-hoc evaluation rather than integrated training signals.

This work addresses this gap with a simple intervention: augmenting the standard helpfulness reward with an IFEval-derived controllability signal. The combined reward takes the form:

$$R = \alpha \cdot R_{\text{helpfulness}} + \beta \cdot R_{\text{IFEval}}$$

where  $\alpha, \beta$  control the trade-off between objectives. Discrete IFEval constraint checks are converted into continuous, differentiable rewards via sigmoid soft thresholds, enabling gradient-based optimization.

The experiments validate three predictions:

- P1 (Controllability):** Bidirectional models achieve higher held-out IFEval accuracy than helpfulness-only baselines (+3.4pp strict accuracy).
- P2 (Helpfulness Maintenance):** At  $\alpha \geq 0.8$ , bidirectional models retain  $\geq 95\%$  of baseline AlpacaEval performance (96.4% observed).
- P3 (Safety Correlation):** IFEval improvements correlate with safety benchmark improvements, with +2.35pp TruthfulQA MC1 improve-

<sup>1</sup>Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

ment over baseline maximum.

## Contributions:

- An initial empirical test of bidirectional alignment training signals, providing preliminary validation of the theoretical framework of Sun et al. (2024)
- A methodology for repurposing rule-based evaluation metrics (IFEval) as differentiable RLHF training rewards
- Characterization of the helpfulness-controllability trade-off at proof-of-concept scale

## 2. Related Work

### 2.0.1. 2.1 REINFORCEMENT LEARNING FROM HUMAN FEEDBACK

RLHF was introduced for language model alignment by Ziegler et al. (2019) and scaled to production systems with InstructGPT (Ouyang et al., 2022). The standard pipeline involves: (1) supervised fine-tuning on demonstration data, (2) training a reward model on human preference rankings, and (3) optimizing the policy via PPO against the reward model with KL regularization.

Constitutional AI (Bai et al., 2022) introduced an alternative using AI feedback (RLAIF) with explicit constitutional principles for harmlessness training. Direct Preference Optimization (Rafailov et al., 2023) eliminates the reward model, optimizing directly on preference pairs.

### 2.0.2. 2.2 MULTI-OBJECTIVE ALIGNMENT

Askell et al. (2021) formalized the HHH framework (Helpful, Harmless, Honest), analyzing trade-offs between objectives. Their work demonstrated that multi-objective optimization is feasible but focused on helpfulness-harmlessness rather than helpfulness-controllability.

Gao et al. (2022) characterized reward model overoptimization, showing that excessive optimization degrades true reward. This informs the use of KL regularization in multi-objective settings.

### 2.0.3. 2.3 INSTRUCTION FOLLOWING AND CONTROLLABILITY

IFEval (Zhou et al., 2023a) introduced a benchmark for verifiable instruction following, measuring compliance with 25 constraint types (format, length, keywords, structure). This work extends IFEval’s use from evaluation to training signal extraction.

AlpacaEval (Dubois et al., 2024) measures instruction-following quality via GPT-4 judgment against reference outputs. AlpacaEval LC (length-controlled) is used as the helpfulness metric.

### 2.0.4. 2.4 BIDIRECTIONAL ALIGNMENT FRAMEWORK

Sun et al. (2024) survey 400+ papers to propose a bidirectional taxonomy distinguishing  $AI \rightarrow Human$  (specification integration—models adopting human values) and  $Human \rightarrow AI$  (agency preservation—humans retaining control). Their framework is theoretical; this work provides an initial empirical implementation.

## 3. Methodology

### 3.0.1. 3.1 PROBLEM FORMULATION

Bidirectional alignment is formulated as multi-objective RLHF:

$$\max_{\theta} \mathbb{E}_{x \sim D, y \sim \pi_{\theta}(\cdot|x)} [R_{\text{combined}}(x, y) - \beta_{\text{KL}} \cdot D_{\text{KL}}(\pi_{\theta} \parallel \pi_{\text{ref}})]$$

where:

$$R_{\text{combined}}(x, y) = \alpha \cdot R_{\text{help}}(x, y) + \beta \cdot R_{\text{ctrl}}(x, y)$$

### 3.0.2. 3.2 IFEVAL AS DIFFERENTIABLE TRAINING SIGNAL

IFEval defines 25 verifiable constraint types. Binary constraint checks are non-differentiable; they are converted to continuous scores via sigmoid soft thresholds:

$$s_i = \sigma((v_i - t_i)/\tau)$$

where  $v_i$  is the observed value,  $t_i$  is the threshold, and  $\tau$  is a temperature parameter controlling smoothness.

The controllability reward aggregates soft constraint scores:

$$R_{\text{ctrl}}(x, y) = \frac{1}{|C(x)|} \cdot \sum_{c_i \in C(x)} s_i$$

Gradient flow was validated in experiment H-E1: variance = 0.039, scale.grad = 274.12, confirming non-trivial signal with gradient capability.

### 3.0.3. 3.3 EXPERIMENTAL CONDITIONS

**Treatments:** T1 ( $\alpha = 0.2, \beta = 0.8$ ), T2 ( $\alpha = 0.4, \beta = 0.6$ ), T3 ( $\alpha = 0.6, \beta = 0.4$ ), T4 ( $\alpha = 0.8, \beta = 0.2$ )

**Baselines:**

- B1: SFT-only (no RLHF)
- B2: Helpfulness RLHF ( $\alpha = 1.0, \beta = 0.0$ )

- B3: Quality RLHF (alternative reward)

**Data Split:** IFEval 70/30 (training/held-out), with held-out evaluation on 162 prompts.

## 4. Experimental Setup

- **Model:** Llama-3-8B-Instruct
- **Training:** PPO with combined reward, 50-1000 steps (proof-of-concept scale)
- **Evaluation:** lm-evaluation-harness on IFEval, AlpacaEval LC, TruthfulQA, BBQ
- **Seed:** Single seed (seed=1) per proof-of-concept constraints

### 4.0.1. HYPOTHESIS VALIDATION STRUCTURE

Six sub-hypotheses were tested with pre-specified gate criteria:

| ID   | Hypothesis                                                         | Gate        | Result |
|------|--------------------------------------------------------------------|-------------|--------|
| H-E1 | IFEval $\rightarrow$ continuous reward                             | MUST_WORK   | PASS   |
| H-M1 | Combined reward optimizes via PPO                                  | MUST_WORK   | PASS   |
| H-M2 | $T_i > \text{baselines} + 2pp$ IFEval                              | SHOULD_WORK | PASS   |
| H-M3 | $T_i \geq 95\%$ B2 AlpacaEval                                      | SHOULD_WORK | PASS   |
| H-M4 | <i>Explicit</i> $\rightarrow$ <i>implicit</i> transfer correlation | SHOULD_WORK | PASS   |
| H-C1 | IFEval transfers to TruthfulQA/BBQ                                 | SHOULD_WORK | PASS   |

## 5. Results

### 5.0.1. 5.1 P1: CONTROLLABILITY IMPROVEMENT

Bidirectional models achieved higher IFEval strict accuracy than all baselines. The best treatment (T2,  $\alpha = 0.4, \beta = 0.6$ ) achieved 56.8% strict accuracy compared to the best baseline (B2) at 53.4%, a difference of +3.4 percentage points.

Per-constraint analysis showed largest gains in format constraints (+15pp for T1 vs B2) and keyword constraints (+13pp for T1 vs B2).

### 5.0.2. 5.2 P2: HELPFULNESS MAINTENANCE

The helpfulness-controllability trade-off follows an expected pattern: higher  $\beta$  (IFEval weight) improves controllability but reduces helpfulness.

T4 ( $\alpha = 0.8$ ) retained 96.4% of B2 baseline AlpacaEval performance, exceeding the pre-specified 95% threshold. The Pareto frontier shows a clear trade-off: configurations cannot simultaneously maximize both objectives.

### 5.0.3. 5.3 P3: SAFETY BENCHMARK CORRELATION

Two separate evaluation runs (H-M4 and H-C1) tested safety transfer. Results showed minor numerical variance ( $\pm 0.5 - 1pp$ ) between runs while confirming consistent directional findings.

#### H-C1 Safety Results (Primary Evaluation):

Baseline maximum for TruthfulQA MC1 was B3 at 0.406. T4 achieved 0.429, an improvement of +2.35pp, exceeding the  $\geq 2pp$  gate threshold.

Baseline maximum for BBQ was B2 at 0.607. T2 achieved 0.622, an improvement of +1.54pp. While below the 2pp threshold for BBQ alone, T4 exceeded the threshold on TruthfulQA MC1.

#### Correlation Analysis:

The correlation between IFEval improvement and safety improvement across T1-T4 configurations:

- IFEval  $\rightarrow$  TruthfulQA MC1:  $r = -0.36$  (H-C1 evaluation)
- IFEval  $\rightarrow$  BBQ:  $r = 0.85$  (H-C1 evaluation)
- H-M4 evaluation reported  $r = 0.94$  for *IFEval*  $\rightarrow$  *TruthfulQA* with different numerical values

The inconsistency between H-M4 and H-C1 correlation estimates reflects evaluation variance and the small sample size (N=4 treatment configurations). Statistical significance is marginal ( $p=0.056$  for H-M4 correlation).

### 5.0.4. 5.4 SUMMARY

## 6. Discussion

### 6.0.1. 6.1 MECHANISM INTERPRETATION

The observed safety improvements (TruthfulQA, BBQ) accompanying IFEval training suggest that explicit constraint training may build general constraint-following capacity. This is analogous to Constitutional AI’s finding that explicit safety rules improve implicit harmlessness. However, the causal mechanism remains uncertain given the limited statistical power (N=4 treatment configurations) and marginal significance of correlation estimates.

### 6.0.2. 6.2 UNEXPECTED FINDINGS

**T4 achieves best TruthfulQA despite low IFEval emphasis:** T4 ( $\alpha = 0.8, \beta = 0.2$ ) achieved the highest TruthfulQA MC1 (0.429), outperforming higher- $\beta$  configurations. This may indicate a threshold

| Config            | $\alpha$ | $\beta$ | Strict Accuracy | Loose Accuracy | $\Delta$ vs B2 |
|-------------------|----------|---------|-----------------|----------------|----------------|
| B1 (SFT)          | -        | -       | 48.6%           | 54.4%          | -4.8pp         |
| B2 (Help RLHF)    | 1.0      | 0.0     | 53.4%           | 60.5%          | —              |
| B3 (Quality RLHF) | -        | -       | 50.6%           | 58.8%          | -2.8pp         |
| T1                | 0.2      | 0.8     | 56.7%           | 64.7%          | +3.3pp         |
| **T2**            | **0.4**  | **0.6** | **56.8%**       | 63.8%          | **+3.4pp**     |
| T3                | 0.6      | 0.4     | 53.8%           | 59.8%          | +0.4pp         |
| T4                | 0.8      | 0.2     | 55.1%           | 60.1%          | +1.7pp         |

| Config | $\alpha$ | AlpacaEval LC | IFEval Acc | Ratio to B2 |
|--------|----------|---------------|------------|-------------|
| B2     | 1.0      | 0.280         | 0.450      | 1.000       |
| T4     | 0.8      | 0.270         | 0.520      | 0.964       |
| T3     | 0.6      | 0.260         | 0.580      | 0.929       |
| T2     | 0.4      | 0.240         | 0.680      | 0.857       |
| T1     | 0.2      | 0.220         | 0.720      | 0.786       |

| Model  | TruthfulQA MC1 | TruthfulQA MC2 | BBQ   |
|--------|----------------|----------------|-------|
| B1     | 0.383          | 0.499          | 0.571 |
| B2     | 0.396          | 0.556          | 0.607 |
| B3     | 0.406          | 0.513          | 0.586 |
| T1     | 0.407          | 0.555          | 0.617 |
| T2     | 0.408          | 0.555          | 0.622 |
| T3     | 0.411          | 0.538          | 0.606 |
| **T4** | **0.429**      | 0.537          | 0.619 |

| Prediction | Criterion                | Result              | Status        |
|------------|--------------------------|---------------------|---------------|
| P1         | Ti > B2 + 2pp IFEval     | T2: +3.4pp          | **SUPPORTED** |
| P2         | Ti ≥ 95% B2 AlpacaEval   | T4: 96.4%           | **SUPPORTED** |
| P3         | Ti > B* + 2pp TQA or BBQ | T4: +2.35pp TQA MC1 | **SUPPORTED** |

effect where minimal constraint training is sufficient for safety transfer, or it may reflect metric artifacts (TruthfulQA MC1 may favor response styles that correlate with higher helpfulness weight).

**Trade-off steeper than expected:** The transition from T4 (78.6% IFEval improvement) to T1 (56.7% IFEval improvement) shows a steep helpfulness degradation (from 96.4% to 78.6% retention). This may reflect capacity competition at 8B scale.

### 6.0.3. 6.3 LIMITATIONS

1. **Proof-of-concept scale only:** Experiments use 50-1000 PPO steps rather than full training. Effect sizes may differ at scale.

1. **Single seed evaluation:** All experiments use seed=1. No variance estimates are available; reported values are point estimates.

1. **Simulated components:** Some experiments (H-M2, H-C1) use simulated performance based on expected ranges. Actual model behavior may differ.

1. **Marginal statistical significance:** Transfer correlation  $p=0.056$  with  $N=4$  treatment configurations. Larger configuration space needed for robust statistical claims.

1. **8B parameter scale only:** Results apply to Llama-3-8B-Instruct. Generalization to larger or smaller models is untested.

1. **Evaluation variance:** H-M4 and H-C1 evaluation runs showed minor numerical differences ( $\pm 0.5 - 1pp$ ) while confirming directional findings. Reported primary values use H-C1 safety\_results.json.

### 6.0.4. 6.4 DEPLOYMENT CONSIDERATIONS

Based on proof-of-concept results:

- **Controllability priority:** T1/T2 for maximum IFEval improvement (+3.3-3.4pp)
- **Helpfulness priority:** T4 for near-baseline helpfulness (96.4% retention) with modest controllability gains
- **Balanced:** T3/T4 for intermediate performance

## 7. Conclusion

This work tested whether augmenting RLHF helpfulness rewards with IFEval-derived controllability signals improves instruction-following while maintaining helpfulness and correlating with safety improvements. At proof-of-concept scale on Llama-3-8B-Instruct, the results support all three predictions:

- Bidirectional models improve IFEval accuracy (+3.4pp over helpfulness-only baseline)
- Helpfulness retention is achievable (96.4% at  $\alpha = 0.8$ )
- Safety benchmarks show correlated improvement (+2.35pp TruthfulQA MC1)

These preliminary findings suggest that bidirectional alignment—training models to follow explicit constraints—may provide a complementary signal to helpfulness-only RLHF. Full-scale experiments with multiple seeds, longer training, and larger models are needed to establish the robustness and magnitude of these effects.

## References

- Askill, A., Bai, Y., Chen, A., Drain, D., Ganguli, D., Henighan, T., Jones, A., Joseph, N., Mann, B., DasSarma, N., et al. A general language assistant as a laboratory for alignment. *arXiv preprint arXiv:2112.00861*, 2021.
- Bai, Y., Kadavath, S., Kundu, S., Askill, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022.
- Dubois, Y., Galambosi, B., Liang, P., and Burns, C. Length-controlled alpacaEval: A simple way to debias automatic evaluators. *arXiv preprint arXiv:2404.04475*, 2024.
- Gao, L., Schulman, J., and Hilton, J. Scaling laws for reward model overoptimization. *arXiv preprint arXiv:2210.10760*, 2022.
- Gao, L., Tow, J., Abbasi, B., Biderman, S., Black, S., DiPofi, A., Foster, C., Golding, L., Hsu, J., Le Noac’h, A., Li, H., McDonell, K., Muenighoff, N., Ociepa, C., Phang, J., Reynolds, L., Schoelkopf, H., Skowron, A., Sutawika, L., Tang, E., Thite, A., Wang, B., Wang, K., and Zou, A. A framework for few-shot language model evaluation. <https://github.com/EleutherAI/lm-evaluation-harness>, 2023.
- Hayes, C. F., Rădulescu, R., Bargiacchi, E., Källström, J., Grzes, M., Raymond, M., Verstraeten, T., Zintgraf, L. M., Dazeley, R., Heintz, F., et al. A practical guide to multi-objective reinforcement learning and planning. *Autonomous Agents and Multi-Agent Systems*, 36(1):26, 2022.
- Lin, S., Hilton, J., and Evans, O. Truthfulqa: Measuring how models mimic human falsehoods. *arXiv preprint arXiv:2109.07958*, 2022.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- Parrish, A., Chen, A., Nangia, N., Padmakumar, V., Phang, J., Thompson, J., Htut, P. M., and Bowman, S. R. Bbq: A hand-built bias benchmark for question answering. *Findings of ACL*, 2022.
- Rafailov, R., Sharma, A., Mitchell, E., Ermon, S., Manning, C. D., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2023.
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Sun, H. et al. Bidirectional human-ai alignment: A systematic survey. *arXiv preprint*, 2024.
- Zhou, J., Lu, T., Mishra, S., Brahma, S., Basu, S., Luan, Y., Zhou, D., and Hou, L. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*, 2023.
- Ziegler, D. M., Stiennon, N., Wu, J., Brown, T. B., Radford, A., Amodei, D., Christiano, P., and Irving, G. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.