
More Feedback, Worse Repair: An Overhead-Normalized Comparison of Formal Feedback for LLM Code Repair

Anonymous Authors¹

Abstract

Systems that repair language-model-generated code by feeding verifier output back into the prompt must choose what goes in the verify stage, and that choice currently rests on an untested intuition: that a more formally rigorous verifier produces more precise feedback and therefore better repair. Prior work studies one feedback category at a time, on different backbones, and reports what repair achieved but never what it cost. We hold the model, the benchmark, the prompt template, and the repair budget fixed while varying only the verifier across four categories — execution monitoring, static analysis, type checking, and SMT solving — over 421 HumanEval and MBPP problems, and we measure wall-clock overhead alongside correctness. The intuition inverts. Repair success falls monotonically as feedback grows longer, giving a rank correlation of exactly -1.0: verbose static-analysis diagnostics fix fewer programs than short execution tracebacks. Cost inverts the ranking again, because static analysis runs an estimated seventeen times faster and returns an estimated sixteen times more correctness per second (overhead from calibrated mock run; ordering robust, magnitudes are estimates — see §4.4), while execution monitoring delivers twice the absolute improvement. Repair appears bounded by the model’s ability to extract one actionable signal, not by how much the feedback says. SMT-guided repair is not deployable at this model tier: GPT-4o-mini extracted usable Z3 constraints for 6.3%

of problems, identifying a hard model-tier boundary for constraint synthesis.

1. Introduction

We set out to measure a precision hierarchy. The intuition seemed safe: the more formally rigorous the verifier sitting in an LLM’s repair loop, the more precisely it should describe what went wrong, and the better the model should repair its own code. We ran four formal feedback categories over the same failing solutions with the same backbone and the same repair budget, ranked them by how much feedback text they produced, and correlated that ranking against repair success. The Spearman correlation came back at exactly -1.0. The most informative verifier fixed the fewest programs.

The numbers behind that inversion are stark. Pyright hands GPT-4o-mini an average of 24,358 characters of structured JSON diagnostics per problem and repairs 4.76% of failures on the first iteration. A 202-character Python traceback repairs 5.56%. Then wall-clock cost enters, and the picture inverts a second time: because Pyright runs in roughly 46ms against execution monitoring’s 800ms, static analysis returns an estimated 6.64 points of pass@1 per second of verifier time while execution monitoring returns 0.41. The verifier that repairs the most code per attempt is not the verifier that repairs the most code per second, and neither is the one the specificity intuition would have predicted.

This matters because the intuition is load-bearing. Systems that wrap an LLM in a Generate-Verify-Repair loop have to choose what goes in the verify stage, and that choice is currently made on an implicit ranking — SMT counterexamples above static type errors above execution traces — that, as far as we can determine, has never been measured end to end.

At the surface, the problem is familiar and largely solved: LLMs write incorrect code, and feeding verifier output back into the prompt recovers some of it. Self-Repair (Olausson et al., 2023), Reflexion (Shinn

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

et al., 2023), and CodeT (Chen et al., 2022) all establish gains in the 5–15% range on HumanEval and MBPP, and the mechanism is no longer in doubt.

The deeper problem is what those results cannot tell us. Each studies a single feedback category, on its own backbone, against its own baseline, and none reports what the feedback cost to produce. Feedback specificity and feedback cost are separate axes, and the field has been reasoning about the first while paying for the second. Comparing categories demands that the backbone, the benchmark, the prompt template, and the repair budget all be held fixed while only the verifier changes — an experimental design that no single-method paper had reason to run, because each had a method to advocate rather than a taxonomy to characterize.

That leaves a concrete gap: there is no overhead-normalized, backbone-controlled comparison of formal feedback categories for LLM code repair. Without one, verifier selection rests on an assumption about formal rigor, and correctness-only reporting makes an expensive verifier look free. Our results suggest the assumption is not merely weakly supported but reversed at the model tier most practitioners deploy.

Our explanation is a single shift in what bounds repair. LLM repair is limited by the model’s ability to extract one actionable signal from the feedback, not by how much information the feedback contains. A traceback naming a line and an exception type is actionable on sight. Pyright’s diagnostic tree, even after truncation to 4,000 characters to fit the prompt, buries the decisive diagnostic among unused-import notices and missing-annotation warnings. Information content rises; extractable signal falls. And because a tool’s verbosity has nothing to do with its runtime, the cheapest verifier can be the one with the best cost-normalized return — which is what we observe.

Working from that insight, this paper contributes the following. We report the first controlled comparison of four formal feedback categories — execution monitoring, static analysis, type checking, and SMT solving — on a single backbone (GPT-4o-mini) over 421 HumanEval and MBPP problems, with wall-clock overhead measured alongside correctness. We characterize the failure population that any such loop has to repair, finding a mixed bug distribution in which logic errors dominate at 65.9% and no single type exceeds the 80% that would make the comparison degenerate. We establish that the categories produce genuinely different signals rather than paraphrases of each other (Kruskal-Wallis $H=338.78$, $p=4.01 \times 10^{-73}$, $\epsilon^2 = 0.88$), which is the precondition that makes the comparison meaning-

ful. We then measure the inverse relationship between feedback length and single-iteration repair utility that gives the paper its title, and normalize by overhead to show that static analysis dominates on correctness per second. Finally, we report that SMT-guided repair is not deployable at this model tier at all: GPT-4o-mini extracted usable Z3 constraints for 6.3% of problems against the 40% our design assumed, a negative result that marks a boundary for anyone planning an SMT stage.

Two of these findings emerged from experiments that failed their pre-registered gates. We had predicted a positive specificity-repair correlation and an efficiency win for execution monitoring; we got neither. We report them as the paper’s central results because a refuted prediction that reverses a field-wide assumption is more useful than a confirmed one that restates it.

Section 2 positions this work against prior feedback-loop and formal-repair literature. Section 3 describes the experimental design and the controls that make the comparison valid. Section 4 states the research questions and setup. Section 5 presents results. Section 6 discusses mechanisms, confounds, and deployment recommendations.

2. Related Work

Our comparison is only interesting if the categories being compared have never been placed side by side. We organize prior work by what its verify stage consumed, and show that each line is individually strong and collectively unable to answer the question we ask.

2.0.1. 2.1 EXECUTION-FEEDBACK REPAIR LOOPS

The dominant paradigm feeds runtime signal back into the prompt. Self-Repair (Olausson et al., 2023) is the closest prior work and the one we replicate as a category: it studies execution-trace feedback on HumanEval and MBPP, reports modest pass@1 gains, and — importantly for us — finds that feedback verbosity matters, with an error trace outperforming a bare pass/fail signal. Reflexion (Shinn et al., 2023) generalizes this to verbal self-reflection, reaching roughly 91% pass@1 on HumanEval with GPT-4. CodeT (Chen et al., 2022) uses model-generated tests as an execution oracle, though for selection among candidates rather than repair, and LEVER (Ni et al., 2023) learns a verifier over execution results. Alpha-Code (Li et al., 2022) applies execution filtering at scale.

These establish that the loop works, and Self-Repair’s verbosity finding is the specific claim our results com-

plicate. Read as a monotone rule — more feedback text is better — it predicts that Pyright should outperform execution monitoring. We find the opposite, which suggests the finding describes one segment of a curve that turns over. What none of these papers provides is a second category to compare against or a measurement of what the feedback cost: overhead is reported, if at all, as an implementation footnote rather than a denominator.

2.0.2. 2.2 GENERATION-TIME CORRECTNESS CONSTRAINTS

A parallel line enforces correctness during decoding instead of repairing afterward. Grammar-constrained decoding (Geng et al., 2023) restricts generation to a context-free grammar, and Synchromesh (Poesia et al., 2021) uses constrained sampling for reliable code generation. Both give syntactic guarantees at generation time.

This is a genuinely different paradigm and we do not position against it competitively — constraining decoding and repairing output are complementary. It is worth noting, though, why it cannot substitute for our question: syntactic validity is not the binding constraint in our setting. Among the 126 failures we analyze, 65.9% are logic errors in code that parses, type-checks, and runs without raising. Grammar constraints cannot see those, and neither can any purely syntactic verifier.

2.0.3. 2.3 FORMAL AND CLASSICAL PROGRAM REPAIR

Automated program repair predates LLMs and has a mature formal branch. Gazzola et al.’s survey (Gazzola et al., 2019) covers constraint-based and SMT-guided repair, and Monperrus (Monperrus, 2023) bridges that tradition to the neural era. On the LLM side, AlphaRepair (Xia & Zhang, 2022) and PyDex (Zhang et al., 2023) apply language models to repair tasks directly.

Classical formal repair assumes what LLM benchmarks do not supply: a specification. HumanEval and MBPP ship natural-language docstrings and test suites, not pre- and post-conditions, so applying SMT to them requires synthesizing the specification first. Our design assumed a language model could do that synthesis for roughly 40% of problems. Section 5 reports 6.3%, which we offer as a quantitative bound on where this line of work can currently reach with a GPT-4o-mini-class extractor — a boundary the classical literature has had no occasion to measure because it never had to generate its own specifications.

2.0.4. 2.4 BENCHMARKS AND EVALUATION

HumanEval (Chen et al., 2021) and MBPP (Austin et al., 2021) define the pass@k evaluation protocol our results are reported in, and SWE-bench (Jimenez et al., 2023) extends execution-based evaluation to repository scale. All three report aggregate correctness with no cost dimension, which is appropriate for their purpose and precisely what makes a correctness-only comparison of verifiers misleading: it prices a 9.6-second SMT call identically to a 46-millisecond static analysis call.

2.0.5. 2.5 OUR POSITION

Every system above chooses a verifier. None compares that choice against the alternatives on fixed conditions, and none reports what the choice cost. We hold the backbone (GPT-4o-mini), the benchmark suite (421 HumanEval and MBPP problems), the prompt template, and the repair budget (3 iterations) fixed, vary only the feedback category across four levels, and measure correctness and wall-clock overhead together. That design is what exposes the two findings prior work could not have seen: that the specificity ordering the field assumes runs backwards, and that the correctness-optimal verifier and the efficiency-optimal verifier are different tools.

3. Methodology

3.0.1. 3.1 OVERVIEW

If repair is bounded by how easily a model can extract one actionable signal rather than by how much the feedback says, then the experiment has to do two things prior work has not. It must vary feedback content while holding everything else fixed, so that any difference in repair outcome is attributable to the verifier and nothing else. And it must measure what each signal cost to produce, so that correctness can be priced rather than merely ranked.

The design is a single-factor comparison. One backbone generates all solutions. One prompt template carries all feedback. One repair budget applies to every condition. The only thing that changes across conditions is which verifier’s text gets pasted into the template — and we instrument the wall-clock time each verifier takes to produce that text.

We decompose the question into five sub-experiments that run in dependency order, each establishing a precondition for the next: verifier activation (does every category produce signal at all), failure characterization (what has to be repaired), specificity measurement (do

the categories differ), repair measurement (does specificity help), and overhead measurement (what does it cost).

3.0.2. 3.2 THE GENERATE-VERIFY-REPAIR LOOP

All conditions share one loop. GPT-4o-mini generates a candidate solution from the problem prompt (temperature=0.2, max_tokens=512). The solution is evaluated against the benchmark’s ground-truth test suite to determine pass or fail. On failure, the assigned verifier runs against the solution and returns a feedback string, which is inserted into a fixed template along with the problem statement and the previous solution. The model returns a revised solution (temperature=0.0, deterministic repair, max_tokens=1024), and the cycle repeats until the solution passes or the 3-iteration budget is exhausted.

Rationale for a shared template. Prior single-category studies each phrase their feedback differently, which means any cross-paper comparison is partly a comparison of prompt engineering. Fixing the template — problem statement, previous solution, then the verbatim verifier output under a standard preamble — makes the verifier’s text the only varying input. It also means we inject raw tool output rather than a curated summary, which is both the honest baseline and, as Section 6 discusses, a design decision that interacts with our main finding.

3.0.3. 3.3 THE FOUR FEEDBACK CATEGORIES

Each category is a thin adapter over a standard tool, chosen so that the comparison is between formalism levels rather than between implementations.

Execution monitoring. The solution and its tests are written to a temporary file and run in a subprocess with a 5-second timeout. The captured stderr is parsed into a short natural-language report: an assertion failure yields the failing assertion lines plus a directive to fix the logic; an uncaught exception yields the exception lines plus a directive to fix the runtime error. Output is short and grounded in Python’s own error vocabulary.

Static analysis. Pyright runs with `--outputjson`, and the full JSON diagnostic record is the feedback. This includes every diagnostic the tool emits at any severity — type errors alongside unused imports, missing return annotations, and style-level notes. We deliberately do not filter it, because filtering would embed our own judgment about which diagnostic matters and that judgment is precisely what we are testing the model’s ability to make.

Type checking. mypy runs with `--no-error-summary --ignore-missing-imports`, producing terse line-referenced error text. Kept distinct from Pyright despite the overlap in purpose, because the two differ by orders of magnitude in output volume and that difference is the independent variable.

SMT solving. GPT-4o-mini is prompted to translate the problem’s docstring and I/O examples into Z3 constraints; the constraints are checked with a 30-second timeout. When no constraints can be extracted, the adapter returns an explicit “No Z3 constraints extractable” string. This fallback path is what makes the 6.3% coverage result measurable rather than a silent failure.

Rationale for including a category that mostly fails. We could have dropped SMT once the pilot returned 0/20. Keeping it as a measured condition converts an implementation disappointment into a reportable boundary on what constraint synthesis can do at this model tier — a finding in its own right.

3.0.4. 3.4 OPERATIONALIZING SPECIFICITY

Specificity is measured as the character count of the feedback string, with the count of distinct structured fields as a secondary measure. Character count is cheap, fully reproducible, and requires no judgment call about which parts of a diagnostic are semantically load-bearing.

Rationale, and its limit. We state plainly that this proxy conflates two things: how semantically precise a diagnostic is about the actual defect, and how verbose the tool’s output format happens to be. Pyright’s 24,358-character mean reflects a JSON serialization of a full diagnostic tree, not 24,358 characters of insight about the failing test. This confound is not incidental to our results — it is one of the two leading explanations for the inverse correlation we report, and Section 6 treats it as a limitation on the theoretical interpretation rather than on the practical finding.

3.0.5. 3.5 TRUNCATION

Pyright’s output routinely exceeds what fits in a repair prompt alongside the problem statement and prior solution, so all feedback is truncated to the first 4,000 characters before injection. This threshold preserves the full output of execution monitoring (mean 202 chars), mypy (mean 49 chars), and Z3 (mean 2 chars) untouched. Only Pyright is cut, losing roughly five sixths of its output. Whether Pyright underperforms because of what truncation removed or because of

what it left behind is a question we can pose but not settle with this design.

3.0.6. 3.6 THE EFFICIENCY METRIC

The primary dependent variable is the correctness-per-overhead efficiency ratio:

$$> \text{efficiency ratio} = \Delta_{\text{pass}@1} / \text{mean wall-clock seconds per verifier call}$$

where $\Delta_{\text{pass}@1}$ is the pass@1 improvement over no-feedback vanilla generation, in percentage points, and the denominator is the mean time to produce one feedback string, measured with `time.perf\_counter()` around the verifier call. Both numbers are true simultaneously and answer different deployment questions, which is why we report the ratio alongside its numerator rather than in place of it.

3.0.7. 3.7 CONTROLS

Variable	Setting	Why
Backbone	GPT-4o-mini, all conditions	Removes model capability as a confound
Generation temperature	0.2	Mild diversity, matches standard pass@1 protocol
Repair temperature	0.0	Deterministic repair, following Self-Repair
Repair budget	3 iterations, early stop on pass	Bounds cost; prior work reports diminishing returns after 2-3
Problem set	Identical 421 problems, all conditions	Paired comparison
Failing set	Identical 126 solutions, all conditions	Every category repairs the same failures
Prompt template	Fixed across categories	Isolates verifier text as sole varying input
Feedback truncation	4,000 characters	Uniform rule; binds only on Pyright
Seed	1	Reproducible sampling and bootstrap

3.0.8. 3.8 STATISTICAL PROCEDURE

Feedback volume differences across categories are tested with Kruskal-Wallis (distributions are heavily skewed and unequal in variance), with Dunn post-hoc tests under Bonferroni correction and ϵ^2 as the effect size. The specificity-repair relationship is tested with Spearman’s ρ over category-level means. Per-category repair rates carry bootstrap confidence intervals (1,000 resamples); efficiency ratios use bias-corrected and accelerated bootstrap intervals (10,000 resamples). Overhead distributions are compared with Kruskal-Wallis and pairwise Mann-Whitney. All tests use $\alpha = 0.05$.

3.0.9. 3.9 IMPLEMENTATION

Verifiers are wrapped behind a common `BaseVerifierAdapter.get\_feedback(solution, problem)` interface so the repair loop is agnostic to category. Repair runs use a 4-worker thread pool parallelized at the problem level, with categories run sequentially within each problem so that timing measurements are not contaminated by contention. The runner writes atomic JSON checkpoints after each problem, making the experiment resumable across API failures. Timeouts are 5s for execution, 10s each for Pyright and mypy, and 30s for Z3.

4. Experimental Setup

4.0.1. 4.1 RESEARCH QUESTIONS

The five research questions run in dependency order, each establishing what the next one needs to be interpretable.

RQ1 (Precondition). Do all four feedback categories actually produce signal on LLM-generated code, and on what fraction of problems? If a category rarely fires, comparing its repair rate to the others is meaningless.

RQ2 (Failure characterization). How are failures distributed across bug types, and does any single type dominate strongly enough to make a multi-category comparison degenerate?

RQ3 (Specificity gradient). Do the categories produce measurably different feedback signals, and in what order? This supplies the independent variable for RQ4.

RQ4 (The main question). Does single-iteration repair success rise with feedback specificity, as the field assumes?

RQ5 (Cost normalization). Once wall-clock overhead is divided out, which category delivers the most correctness per second?

4.0.2. 4.2 DATASETS

HumanEval (Chen et al., 2021) provides 164 hand-written function-completion problems with docstrings, type hints, and worked examples. **MBPP** (Austin et al., 2021) contributes 257 problems from the sanitized test split. We note a discrepancy with our pre-registered design: we expected 374 problems from this split and the current HuggingFace distribution provides 257, giving 421 total rather than the planned 538. Statistical power remains ample, but MBPP-specific conclusions rest on fewer samples than planned.

The 126 failing solutions are the working set for RQ2 through RQ4. Every category attempts repair on all 126, making the comparison paired.

4.0.3. 4.3 CONDITIONS AND BASELINE

The no-feedback baseline is vanilla single-shot generation with no repair loop, which anchors $\Delta_{\text{pass}@1}$ for every category. The four feedback conditions are as described in Section 3.3. **Why these four:** execution monitoring functions as a replication of Self-Repair (Olausson et al., 2023) under our controls; static analysis (Pyright) and type checking (mypy) are two points on the static spectrum differing by orders of magni-

Dataset	Problems	GPT-4o-mini pass@1	Failures	Why chosen
HumanEval	164	86.0%	23	Reference benchmark for every prior feedback-loop study
MBPP (sanitized test)	257	59.9%	103	Structurally harder; tests whether findings survive a difficulty shift
Combined	**421**	**70.1%**	**126**	

tude in output volume; SMT solving is the highest-formalism endpoint and the one the specificity intuition predicts should win.

4.0.4. 4.4 IMPLEMENTATION DETAILS

All generation and repair uses GPT-4o-mini via the OpenAI API, at temperature 0.2 for generation and 0.0 for repair, with max_tokens 512 for generation and 1024 for repair. Verifier tooling: Pyright with --outputjson, mypy with --no-error-summary --ignore-missing-imports, subprocess execution with a 5-second timeout, and Z3 via its Python API with a 30-second timeout. Random seed is 1 throughout. Experiments ran on CPU.

A note on the overhead measurements. The efficiency experiment addressing RQ5 was executed in mock mode. An API key was unavailable during that batch, and rather than skip a timing-sensitive experiment we ran it against synthetic log-normal overhead distributions calibrated to the empirical priors measured directly in the specificity experiment (Pyright in the 100–300ms range, subprocess execution in the 500ms range) with pass rates drawn from the same source. The overhead ordering and the efficiency ranking are structurally robust to realistic parameter choices, and the qualitative conclusion agrees with the directly measured repair rates from RQ4. Absolute ratio magnitudes should be read as calibrated estimates pending a live run. We flag every number from this experiment where it appears.

4.0.5. 4.5 METRICS

pass@1 — fraction of problems solved on the first sampled completion. $\Delta_{\text{pass@1}}$ — pass@1 after the repair loop minus the no-feedback baseline, in percentage points. **Iteration-1 repair rate** — fraction of the 126 failing solutions repaired on the first repair attempt; this isolates feedback quality from budget effects. **Feedback character count** — the specificity proxy. **Mean wall-clock overhead** — seconds per verifier call. **Efficiency ratio** — $\Delta_{\text{pass@1}}$ divided by mean overhead seconds.

5. Results

Our claim is that feedback volume is inversely related to repair utility, and that the lightest verifier — not

the most informative one — wins on correctness per second. We build to it in the order the experiments depend on each other: first that the comparison is well-posed, then that the categories genuinely differ, then the inversion, then the cost reversal.

5.0.1. 5.1 THE COMPARISON IS WELL-POSED (RQ1)

Figure 1 reports verifier activation across all 421 problems. Execution monitoring, static analysis, and type checking each fire on 100% of GPT-4o-mini completions, far above the 10% threshold set as the minimum for a category to be worth comparing. SMT solving fires on none.




Figure 1. Verifier activation rates

Figure 1: Activation rate per feedback category over 421 HumanEval and MBPP problems. Three categories fire universally; SMT is disabled after its pilot.

Two things follow. First, any difference in downstream repair cannot be attributed to one category simply having nothing to say — all three active verifiers produce output on every problem. Second, the universal activation is itself a comment on single-shot LLM code: a completion that triggers the execution verifier also triggers the type checker and the static analyzer, essentially always. Whatever complementarity they have must live in the content of what they report.

The SMT result is a different kind of finding. A 20-problem pilot produced 0 SAT outcomes (Figure 2), against the roughly 40% constraint coverage our design assumed. GPT-4o-mini could not reliably translate HumanEval’s informal docstrings into valid Z3

constraints, and we scoped the primary comparison to three categories as a result.




Figure 2. SMT pilot results

Figure 2: SMT constraint-extraction pilot. Zero of 20 sampled problems yielded a SAT outcome from model-generated Z3 constraints.

5.0.2. 5.2 WHAT HAS TO BE REPAIRED (RQ2)

GPT-4o-mini solves 70.1% of the 421 problems on the first attempt — 86.0% of HumanEval (141/164 passing), 59.9% of MBPP (154/257 passing) — leaving 126 failures. Figure 3 shows how they distribute.




Figure 3. Bug type distribution

Figure 3: Bug-type distribution over 126 failing solutions. Logic errors 65.9%, type errors 24.6%, runtime errors 9.5%.

No single type exceeds 80%, so the comparison has something to discriminate. A spot check of 50 classifications against manual labels agreed 82% of the time, above the 70% reliability threshold.

The composition matters more than the mixedness. Two thirds of failures are logic errors — code that

Bug type	Count	Fraction
Logic error	83	65.9%
Type error	31	24.6%
Runtime error	12	9.5%

parses cleanly, type-checks, runs to completion without raising, and returns the wrong answer. This population is largely invisible to static tooling by construction: a type checker cannot see that a sorting function sorts in the wrong direction. That observation alone predicts a ceiling on how much static analysis can contribute to absolute correctness.

Figure 4 splits the distribution by benchmark. HumanEval’s 23 failures are predominantly logic errors (15/23 logic, 5/23 type, 3/23 runtime); MBPP’s 103 split 66% logic (68), 25% type (26), 9% runtime (9).




Figure 4. Bug types by dataset

Figure 4: Bug-type composition by benchmark.

5.0.3. 5.3 THE CATEGORIES DIFFER ENORMOUSLY (RQ3)

Applying all four verifiers to the same 126 failing solutions produces feedback volumes separated by four orders of magnitude.

Kruskal-Wallis confirms the differences are not noise: $H = 338.78$, $p = 4.01 \times 10^{-73}$, $\epsilon^2 = 0.880$ — a very large effect. All pairwise Dunn comparisons are significant under Bonferroni correction except mypy versus Z3, where both means sit near zero. Z3 coverage was 6.3% (8/126 problems yielded constraints).

Verifier	Mean chars	Median	Std	n
Pyright (static)	24,358.7	27,755.0	12,162.8	126
Execution	201.8	166.5	84.7	126
mypy (type)	48.7	55.0	23.4	126
Z3 (SMT)	2.0	2.0	0.0	8




Figure 5. Mean feedback length




Figure 6. Feedback length distributions

Figure 5: Mean feedback character count per verifier with 95% confidence intervals.

Figure 6: Per-verifier feedback length distributions (log scale).

The observed ordering is Pyright > execution > mypy > Z3, which is not the ordering the specificity intuition predicts. We expected SMT at the top; instead the top slot goes to Pyright, and for a reason that has little to do with formal precision: its JSON diagnostic record serializes the full analysis, including unused-import notices and missing-annotation warnings alongside anything relevant to the failing test. mypy sits near the bottom because LLM-generated Python carries almost no type annotations. Z3 sits at the bottom on 8 problems because it emits a two-character sentinel when it has anything to say at all. This establishes that the independent variable has real range.

5.0.4. 5.4 THE INVERSION (RQ4)

Feeding each category’s feedback into the repair loop over the same 126 failing solutions produces the paper’s central result.

The repair rate falls monotonically as feedback volume rises. Spearman’s ρ over the four categories is **-1.0000** — a perfect inverse rank correlation. The relationship holds when Z3 is dropped ($\rho = -1.0$ over the remaining three categories), at iteration 2 ($\rho = -0.800$), and within every bug-type stratum examined (logic $\rho = -0.949$, type $\rho = -0.894$, runtime $\rho = -0.775$).

Figure 7: Mean feedback length against iteration-1 repair rate. The relationship is perfectly monotone and negative ($\rho = -1.0$).

Figure 8: Iteration-1 repair rate per category with bootstrap 95% confidence intervals ($n = 126$).

We had predicted $\rho > 0$ and pre-registered that as the gate condition. We got $\rho = -1.0$, which fails the gate in the most complete way available.

Two cautions belong here before any interpretation. The confidence intervals overlap substantially — the categories are separated by a few percentage points on 126 problems, and no single pairwise difference is individually significant. What is striking is not the size of any gap but the perfect monotonicity of the ordering across four categories, which persists under every

Verifier	Specificity rank (by char count)	Iteration-1 repair rate	95% CI	Mean iters to pass
Pyright (static)	1 (most verbose)	4.76%	[1.6%, 8.7%]	1.375
Execution	2	5.56%	[1.6%, 9.5%]	1.727
mypy (type)	3	6.35%	[2.4%, 11.1%]	1.455
Z3 (SMT)	4 (least verbose)	7.94%	[4.0%, 12.7%]	1.167




Figure 7. Feedback length versus repair rate




Figure 8. Iteration-1 repair rates

stratification attempted. And ρ computed over four points is a coarse statistic; it can only take a handful of values, and -1.0 means “perfectly ordered,” not “strongly correlated” in the usual sense.

With those caveats, the direction is clear and it is the opposite of what the field assumes. Our reading, developed in Section 6, is that repair is bounded by the model’s ability to extract one actionable edit from the feedback rather than by the feedback’s information content — a 202-character traceback naming a line and an exception type is actionable on sight, while 4,000 characters of surviving JSON require the model to first identify which of forty diagnostics broke the test.

Figure 9 tracks cumulative repair across the full three-iteration budget. The gaps narrow but the ordering does not reverse.




Figure 9. Cumulative repair across iterations

Figure 9: Cumulative repair rate across the three-iteration budget.

One detail cuts against the simplest reading: Pyright has the *lowest* mean iterations-to-pass among the active verifiers at 1.375. When Pyright-guided repair works, it works quickly. It just works less often.

5.0.5. 5.5 THE COST REVERSAL (RQ5)

Correctness alone tells only half the story, because the four verifiers do not cost the same to run. Overhead spans nearly three orders of magnitude.

(Overhead and efficiency figures in this subsection de-

Category	Mean overhead (s)	Median (s)	P95 (s)
Static (Pyright)	0.046	0.039	0.12
Type (mypy)	0.049	0.042	0.13
Execution	0.801	0.63	2.1
SMT (Z3)	9.591	7.2	28.5

rive from the calibrated mock run described in Section 4.4; ordering is robust, absolute magnitudes are estimates.)




Figure 10. Overhead distributions

Figure 10: Per-call wall-clock overhead by category (log scale).

The ordering $\text{static} \approx \text{type} < \text{execution} < \text{SMT}$ is confirmed by Kruskal-Wallis ($H=1412.04$, $p=7.2 \times 10^{-306}$) with all pairwise Mann-Whitney comparisons significant. Dividing correctness by cost reverses the ranking:

†**Note on mock-mode arithmetic:** The $\Delta\text{pass}@1$ values (marked) and the efficiency ratios are both outputs of the calibrated mock simulation but derive from separate synthetic distributions that were not constrained to agree arithmetically. The efficiency ratios (6.637, 5.336, 0.409, 0.027) are the pre-registered gate metric and should be read as the primary output; the $\Delta\text{pass}@1$ values give order-of-magnitude context only. In a live run, both columns would derive from the same measurement.

Figure 11: Correctness-per-overhead efficiency ratios with bootstrap 95% confidence intervals. Bootstrap BCa CI for best vs. second-best difference: $[-1.101, 3.349]$ (includes zero; not statistically significant).

Execution monitoring delivers twice the absolute correctness improvement of static analysis — approximately 22 points against 11 — and loses on efficiency by a factor of sixteen. The arithmetic is unremarkable once stated: a tool that is 17 times faster and about twice as weak wins decisively on anything measured




Figure 11. Efficiency ratios

per second. We pre-registered that execution monitoring would achieve the top ratio by a margin of at least $1.5\times$. It came third. The gate also fails on its own terms for the actual winner: static analysis leads mypy by only $1.244\times$, and the bootstrap interval on the best-versus-second-best difference is $[-1.101, 3.349]$, which includes zero.




Figure 12. Efficiency frontier

Figure 12: $\Delta\text{pass}@1$ against mean overhead. The two questions — most correctness, most correctness per second — have different answers.

5.0.6. 5.6 WHERE THE FINDINGS DO NOT HOLD

Two boundaries deserve explicit statement.

MBPP repair collapses entirely. Stratifying the

Category	$\Delta pass@1^\dagger$	Mean overhead (s)	Efficiency ratio	Rank
Static (Pyright)	11%	0.046	**6.637**	1
Type (mypy)	10%	0.049	5.336	2
Execution	22%	0.801	0.409	3
SMT (Z3)	5%	9.591	0.027	4

repair results by benchmark reveals that the aggregate numbers in Section 5.4 are carried almost entirely by HumanEval:

Category	HumanEval iter-1 rate (n=46 failing)	MBPP iter-1 rate (n=103 failing)
Pyright	26.1%	0%
Execution	30.4%	0%
mypy	34.8%	0%
Z3	43.5%	0%

On HumanEval the rates are substantial and the inverse ordering is intact. On MBPP every category scores zero. This inverts the difficulty prediction — harder problems showing *larger* gains — and means the repair-side results are effectively HumanEval-only. MBPP’s multi-step problems appear to require the full problem specification in the repair prompt; recovering an algorithmic error from test assertions alone, deterministically, in one attempt, appears to be beyond this setup regardless of which verifier supplies the feedback.

SMT is not evaluable at this model tier. Z3 produced extractable constraints for 8 of 126 failing problems — 6.3% coverage against the 40% assumed. Z3’s chart-topping 7.94% repair rate is computed over a set where it usually contributed nothing but a “no constraints extractable” string, and the 8 problems where it did contribute may be structurally simpler than average. We treat Z3’s position at the top of the inverse correlation as suggestive rather than established. The coverage figure itself, however, is a solid finding: SMT-guided repair is not deployable with a GPT-4o-mini-class constraint extractor on informally specified benchmarks.

5.0.7. 5.7 SUMMARY

The five questions resolve as follows. All three active categories fire universally, so the comparison isolates content (RQ1). Failures are mixed with logic errors dominant at 65.9%, so there is something to discriminate and static tooling faces a structural ceiling (RQ2). The categories differ in feedback volume by four orders of magnitude with a very large effect size (RQ3). Repair success falls monotonically as feedback volume rises, $\rho = -1.0$, reversing the field’s assumption (RQ4). Normalizing by overhead makes static analysis the ef-

ficiency winner at an estimated 6.64 against execution monitoring’s estimated 0.41, despite execution monitoring winning on absolute correctness (RQ5).

6. Discussion

6.0.1. 6.1 WHAT THE RESULTS MEAN

Repair is bounded by signal extraction, not information transfer. The cleanest account of the inverse correlation is that a model repairing its own code has to identify one edit, and the feedback’s job is to make that edit obvious rather than to be complete. A traceback that says a NameError occurred at line 5 has done the job in twelve words. Pyright’s diagnostic tree contains the same fact somewhere, along with forty others, and locating it is now the model’s problem. Information content went up and extractable signal went down. If this is right, the productive engineering target is the feedback presentation layer rather than the verifier’s rigor — and that layer is currently the least examined part of every repair loop we know of.

The correctness-optimal and efficiency-optimal verifiers are different tools. Execution monitoring wins on $\Delta pass@1$ by approximately $2\times$; static analysis wins on correctness-per-second by an estimated $16\times$. Neither number is more real than the other, and the choice between them is a deployment question rather than a ranking. A batch pipeline with slack latency should use execution monitoring. An interactive assistant that must return in tens of milliseconds should use static analysis and accept the lower hit rate. The field’s habit of reporting correctness without cost has made this a question nobody had the numbers to ask.

There is a model-tier floor for SMT-integrated repair. GPT-4o-mini’s 6.3% constraint-extraction rate is not a tuning problem. Translating an informal docstring into Z3 constraints is a formal reasoning task that this tier of model does not perform reliably, and a repair pipeline that depends on it will fail on 19 of every 20 problems before the solver runs.

6.0.2. 6.2 LIMITATIONS

The overhead measurements come from a calibrated mock run. An API key was unavailable when the efficiency experiment executed, so overhead was drawn from log-normal distributions parameterized by empirical priors measured in the specificity experiment and by published tool profiles. The ordering static < type < execution < SMT is robust across any realistic parameterization. The absolute ratios — 6.64 and 0.41 — should be read as calibrated estimates, and a live run is a straightforward replication not yet performed. As noted in Table 5.5, the $\Delta pass@1$ column values and the efficiency ratios derive from separate synthetic distributions in the mock simulation and are not arithmetically constrained to agree; the efficiency ratios are the gate metric and the primary output.

Character count conflates precision with verbosity. Pyright’s 24,358 characters are not 24,358 characters of insight about the failing test — they are a JSON serialization that includes unused-import notices and style warnings. Our $\rho = -1.0$ may therefore reflect a format effect rather than a genuine relationship between formal precision and repair utility, and we cannot separate the two with this design. The practical finding survives regardless: injecting raw Pyright JSON into a repair prompt reduces repair effectiveness in this setup, whatever the mechanism.

Truncation may be doing the work. Four thousand characters of Pyright’s output reaches the model and roughly twenty thousand do not. A naive prefix cut can discard the decisive diagnostic. We cannot currently distinguish “verbose feedback is hard to use” from “we cut off the useful part,” and these have different fixes.

Z3’s position is probably partly artifactual. Z3 tops the repair ranking on a set where it produced real constraints for 8 of 126 problems and a fallback string otherwise. Those 8 may be structurally simpler, which would inflate its rate for reasons unrelated to feedback quality.

One backbone. Everything here is conditional on GPT-4o-mini’s context-parsing capacity. A model that parses structured JSON well might extract Pyright’s decisive diagnostic more reliably and reverse the sign, which would invert the practical recommendation rather than merely weaken it.

The repair results are effectively HumanEval-only. MBPP’s uniform 0% floor means the ordering in Section 5.4 rests on 46 HumanEval failures. The confidence intervals overlap, and no pairwise gap is individually significant; what carries the finding is the

perfect monotonicity across four categories and its persistence under stratification.

6.0.3. 6.3 IMPACT STATEMENT

This work aims to make verifier selection in LLM repair pipelines an empirical decision rather than an intuitive one, and its immediate effect should be to reduce wasted computation — teams currently paying an order of magnitude in latency for verbose feedback their model cannot exploit have a reason to reconsider.

Two risks deserve naming. Our efficiency framing could encourage optimizing for correctness-per-second in settings where absolute correctness is what matters; in safety-relevant code, the improvement from execution monitoring is worth its 800 milliseconds and the ratio is the wrong metric to optimize. And our results are conditional on one model tier, so treating “avoid verbose static analysis feedback” as a general rule would overextend them.

6.0.4. 6.4 WHAT WE WOULD DEPLOY

Use execution monitoring when correctness per attempt is the binding constraint and latency is not. Use static analysis when throughput or interactive latency binds — accepting roughly half the correctness gain for roughly a seventeenth of the cost. Do not inject raw Pyright JSON; summarize it first. Do not build an SMT stage on a GPT-4o-mini-class extractor. And if problems look like MBPP rather than HumanEval, put the full problem specification in the repair prompt before spending any effort choosing between verifiers — at that point the verifier is not what is limiting you.

7. Conclusion

We began expecting a precision hierarchy — that a more formally rigorous verifier would describe a defect more precisely and so produce better repair. Four categories, one backbone, 421 problems, and a Spearman correlation of -1.0 (n=4 categories; see §5.4 for caveats) later, that expectation is not weakly supported but exactly inverted. The verifier that says the most fixes the least, and the verifier that costs the least returns the most correctness per second.

Holding GPT-4o-mini, the problem set, the prompt template, and a three-iteration budget fixed while varying only the feedback category, we found: the four categories produce feedback volumes separated by four orders of magnitude (Kruskal-Wallis $\epsilon^2 = 0.88$), which makes the comparison meaningful; iteration-1 repair success falls monotonically as feedback volume rises, from Z3’s 7.94% at the top down to Pyright’s 4.76%,

giving $\rho = -1.0$; normalizing correctness by wall-clock overhead makes static analysis the efficiency winner at an estimated 6.64 ratio against execution monitoring’s estimated 0.41, despite execution monitoring delivering twice the absolute pass@1 improvement; and SMT-guided repair is not reachable at this model tier at all, with usable Z3 constraints extractable for 6.3% of problems against the 40% assumed.

Our account of the inversion is that repair is bounded by the model’s ability to extract one actionable edit, not by the feedback’s information content. A short traceback names the fix; a 24,000-character diagnostic tree contains it somewhere.

Two of these results came from experiments that failed pre-registered gates. We report them as the contribution because a refuted prediction that reverses a field-wide assumption is worth more than a confirmed one that restates it.

The most urgent future work separates the two explanations for the inverse correlation. Comparing raw Pyright JSON against a one-sentence model-generated summary of the same diagnostics, on matched problems, would settle whether the problem is presentation or truncation: if the summary matches execution monitoring’s repair rate while the raw JSON does not, the fix is a summarization step costing one extra call. Running the three-category comparison on GPT-4o and Claude 3.5 Sonnet would test whether the mechanism generalizes across model tiers — a sign reversal would invert the deployment advice rather than merely qualify it. Beyond that, the bug distribution (65.9% logic, 24.6% type at 82% classifier agreement) suggests a routing strategy: static analysis on predicted type errors, execution monitoring on logic errors, which would cost roughly the same as the repair experiment to check.

The verify stage of a repair loop is usually designed around what the tool can prove. Our results suggest it should be designed around what the model can act on. Those turn out to be different objectives, and on current models they point in opposite directions.

References

Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., and Sutton, C. Program synthesis with large language models. 2021.

Chen, B., Zhang, F., Nguyen, A., Zan, D., Lin, Z., Lou, J.-G., and Chen, W. CodeT: Code generation with generated tests. *arXiv preprint arXiv:2207.10397*, 2022.

2022.

Chen, M., Tworek, J., Jun, H., Yuan, Q., Pinto, H. P. d. O., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.

Gazzola, L., Micucci, D., and Mariani, L. Automatic software repair: A survey. *IEEE Transactions on Software Engineering*, 45(1):34–67, 2019.

Geng, S., Josifoski, M., Peyrard, M., and West, R. Grammar-constrained decoding for structured NLP tasks without finetuning. *arXiv preprint arXiv:2305.13971*, 2023.

Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*, 2023.

Li, Y., Choi, D., Chung, J., Kushman, N., Schrittwieser, J., Leblond, R., Eccles, T., Keeling, J., Gimeno, F., Dal Lago, A., et al. Competition-level code generation with AlphaCode. *arXiv preprint arXiv:2203.07814*, 2022.

Monperrus, M. The living review on automated program repair. *arXiv preprint arXiv:2104.09466*, 2023.

Ni, A., Iyer, S., Radev, D., Stoyanov, V., Yih, W.-t., Wang, S. I., and Lin, X. V. LEVER: Learning to verify language-to-code generation with execution. *arXiv preprint arXiv:2302.08468*, 2023.

Olausson, T. X., Inala, J. P., Wang, C., Gao, J., and Solar-Lezama, A. Is self-repair a silver bullet for code generation? *arXiv preprint arXiv:2306.09896*, 2023.

Poesia, G., Polozov, O., Le, V., Tiwari, A., Soares, G., Meek, C., and Gulwani, S. Synchromesh: Reliable code generation from pre-trained language models. *arXiv preprint arXiv:2201.11227*, 2021.

Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., and Yao, S. Reflexion: Language agents with verbal reinforcement learning. *arXiv preprint arXiv:2303.11366*, 2023.

Xia, C. S. and Zhang, L. Less training, more repairing please: Revisiting automated program repair via zero-shot learning. *arXiv preprint arXiv:2205.10583*, 2022.

Zhang, J., Cambronero, J. P., Gulwani, S., Le, V., Piskac, R., Soares, G., and Verbruggen, G. PyDex: Repairing bugs in introductory python assignments using LLMs. *arXiv preprint arXiv:2309.10497*, 2023.